

# AWS Lambda – Function as a Service

## Serverless Compute

Content Prepared By: Chandra Lingam, Cotton Cola Designs LLC

For Distribution With AWS Certification Course Only

Copyright © 2017 Cotton Cola Designs LLC. All Rights Reserved.

All other registered trademarks and/or copyright material are of their respective owners



# Serverless Compute – What is it?

Application still runs on “servers”

All server management is done by AWS

## Servers to Serverless

EC2 – You manage the Server

RDS – You pick server configuration and AWS manages

Lambda, SQS, S3 .. – Serverless. Fully managed by AWS

# Benefits of Lambda

No Servers to manage – just write your code and upload

Continuous Scaling – Lambda automatically scales by running your function in parallel

No charge when code is not running - don't pay for idle infrastructure

Run code in response to a trigger

# Lambda Language Support

Node.js

Python

Java

C# (.NET Core)

AWS Manages Servers and Execution Environment

Cannot customize OS or Execution environment

# Lambda Invocation – Usage Scenarios

Push: S3 to Lambda

Polling: Kinesis Streams to Lambda

Invoke: App to Lambda

Event Source Mapping – Where is it maintained?

Supported Event Sources & Sample Event Data

# IAM Permissions

- Lambda Function needs permission to access resources
  - Granted through Execution Role – attached to the function
  - Lambda Assumes Roles when executing the function
- Who can invoke the Lambda Function?
  - Resource Based Policies at Function Level
  - IAM Roles, User, Group Policies

# Demo – Hello World

## Example - 1

- Returns Hard Corded String

## Example – 2

- Returns back the event data that was sent
- Write, Configure, Test and check CloudWatch Logs
- Invocation Example from Console and CLI
- Async and Sync Invocation

# Demo – API Gateway Example

- Expose Lambda function as RESTful API
- Configure API Gateway



# Demo – S3 to Lambda

- Push Example
- Trigger Lambda Code on object creation
- Lambda reads the object and prints metadata

# Demo – SNS to Lambda

- Subscribe Lambda to SNS Topic
- Process Events
- Push example

# Demo – Lambda to SQS

- Polling Example
- Scheduled invocation of Lambda function using CloudWatch Event Scheduler
- Polls configured SQS queue
- Fan out – Parallelize processing of messages
- Use Environment Variable to store the Queue URL

# Demo – Version and Alias

- Create Multiple Versions of Function
- Latest version invocation
- Specific version invocation
- Alias - Hide version from Event Sources
- Update Alias to point to new versions

# AWS Lambda - Containers

- Packages your code and launches a [container](#)
- Takes time to setup container and bootstrap – possibly for every lambda invocation
- AWS Lambda tries to reuse containers for subsequent invocations
- [Under Reuse:](#)
  - Any declarations remains initialized – for example: database connections
  - Scratch space /tmp is available – can use it to cache
  - Background process or callback that were not completed may resume!

# Function Configuration

- [Compute Resources](#) Needed – Specify total memory needed. Lambda allocates CPU capacity in proportion
  - Increments of 64 MB
- Maximum Execution Timeout - Configurable up to max 300 seconds (5 minutes)
- IAM Execution Role – Grants necessary permissions for function to access other resources
- Handler Name – Main method where Lambda Execution begins

# Invocation Types

- Synchronous (RequestResponse)
- Asynchronous (Event)

You can specify invocation type through code or CLI

When other AWS Services invoke Lambda, invocation type is predetermined for each of these services

# Invocation Types – From AWS Services

Some examples...

S3 – Asynchronous invocation

Cognito – Synchronous invocation

Kinesis/DynamoDB Streams – Lambda polls for data and invokes function Synchronously



# Concurrent Executions – Non Stream based

- Each published event is a unit of work
- Number of events dictate concurrency
- $\text{Concurrency} = \text{Number of events per second} * \text{function duration}$

## Example

Function duration = 3 seconds

Events per second = 10

Concurrent Executions =  $10 * 3 = 30$

- 1000 concurrent execution limit across all your functions

# Concurrent Executions - Stream based

Number of shards per stream is the unit of concurrency.

Example: 100 shards would trigger 100 lambda functions running concurrently

One lambda function processes events on a shard in the order that they arrive

Kinesis Streams and DynamoDB Streams are natively supported by Lambda

# Throttling & Error Handling

- Default limit = 1,000 concurrent executions (soft limit)
- Any request rate above the limit is throttled
- Stream based
  - Keeps trying until data expires (up to 7 days for Kinesis)
  - Throttled events in a Shard are blocking – strictly processed in-order
  - Non-throttled shard can continue processing

# Throttling & Errors – Non Stream Based

- Synchronous (RequestResponse Type)
  - Invoking application receives 429 error
  - Application is responsible for retries
  - AWS Event Sources may have their own retries
- Asynchronous (Event Type)
  - Events are queued
  - AWS Lambda automatically retries for up to six hours with delays

# Authoring Functions

Tools Available for Authoring

General Structure of a Lambda function

Deployment Package

# Version and Alias

[Manage Version and Alias](#)

# Environment Variables

Pass settings to Lambda Function using [Environment Variables](#)

# VPC Support

- Lambda by default runs inside system-managed VPC
- Lambda cannot access resources in a customer VPC
- Function can be configured to run inside customer VPC
  - VPC
  - Subnet(s)
  - Security Group(s) to use
  - Lambda attaches Elastic Network Interface (ENI) based on above configuration to connect to resources
  - ENI can result in [additional startup delay](#)
- Ensure VPC has [sufficient](#) addresses in subnet and ENI



# VPC Support

- Lambda on system-managed VPC can access internet directly
- Lambda on Customer VPC for outbound internet access needs NAT device
  - Attached ENI has only private address
  - Configure Lambda to run in private subnets
  - Ensure NAT instance or Gateway is configured for VPC
  - Required for accessing other AWS Services

# Monitoring

Monitoring Lambda Functions

Metrics

# Dead Letter Queues

- Async – failed invocation is retried twice and then discarded
- Configure failed invocations to be sent to SQS queue or SNS topic

# X-RAY

Trace Lambda Calls with [AWS X-Ray](#)

# AWS Lambda Limits

[Table: Limits](#)

# Pricing

[Table: Pricing](#)

# Recommended Videos

[AWS re:Invent 2016: Serverless Architectural Patterns and Best Practices \(ARC402\)](#)