

# Cloud Databases

## Fully Managed Database Services

Content Prepared By: Chandra Lingam, Cloud Wave LLC

For Distribution With AWS Certification Course Only

Copyright © 2018 Cloud Wave LLC. All Rights Reserved.

All other registered trademarks and/or copyright material are of their respective owners



# AWS Cloud Database Offerings

- [Relational](#) Database Service (RDS)
- [DynamoDB](#) NoSQL Database Service
- [Redshift](#) Data warehouse Service
- [ElastiCache](#) in-memory cache Service
- Database [Migration](#) Service

# Relational Databases

- Data Stored as Tables, Columns and Relations
- Rigid Schema
- Usually normalized to minimize duplication
- Very popular with widespread usage
- SQL Query Language
- Scale-up processing

# Movie Example

*Movie Table*

ID	Title	Year	Date Released
1	Rush	2013	2013-09-02
2	Prisoners	2013	2013-08-30

*Movie\_Actor  
Table*

Movie ID	Actor ID
1	100
1	101

*Movie\_Genre  
Table*

Movie ID	Genre ID
1	1
1	3

*Actor Table*

ID	Actor Name
100	Daniel Bruhl
101	Chris Hemsworth

*Genre Table*

ID	Genre
1	Action
2	Biography
3	Drama

Complex Query, Transaction,  
Typically used for “Write” heavy  
data

# Movie – Denormalized Example

Movie Table

ID	Title	Year	Date Released	Actor_1	Actor_2	Genre_1	Genre_2
1	Rush	2013	2013-09-02	Daniel B..	Chris H..	Action	Drama
2	Prisoners	2013	2013-08-30	Hugh J..	Viola D..	Crime	Drama

Easy to Query, More Storage, Duplication, Typically used for “Read-Only” Data

# NoSQL Databases

- High Performance, Non-relational databases
- Variety of data storage
  - Key-Value Pairs
  - Document Store
  - Graph
  - Column Oriented
- Denormalized, Easy to change “Schema”
- Some support SQL Query Language
- Scale-out processing

# NoSQL Document Movie Example

Review: OneMovie.json Example File

Reality: NoSQL often “normalized” into separate tables to optimize for specific use case. For example: Movie in one table, Movie Reviews in another table, Ratings in third table and so forth

# Columnar Databases

- Optimized for reading and writing specific columns
- Ideal for analytic query performance – when analyzing select columns
- Reduced Disk I/O requirement
- High level of compression
- Scale-out processing
- [Example: Redshift Columnar Storage](#)



# In-Memory Storage

- In-memory database optimized for read-heavy workloads and compute intensive workloads
- Extremely fast sub-millisecond response time
- Reduce traffic hitting backend databases
- Ideal for social networking, gaming leadership boards, API request limit throttling, recommendation engines

# AWS ElasticSearch Service

- Text Search Engine
- Faceting Search allows categorization of search results
  - Price Range
  - Product Types (Video, Book, Toys)
  - Brands
- Suggestions as you type
- Reduce traffic hitting backend databases
- Very fast response times
- Used for application search, log analytics, application monitoring, clickstream analytics

# Relational Versus NoSQL

[Table: Comparison Between Relational and NoSQL](#)

[Terminology: Relational and NoSQL](#)

# Relational Database Service (RDS)

# Relational Database Service

- Six Popular Database Engines: Aurora, MySQL, MariaDB, PostgreSQL, Oracle, SQL Server
- RDS handles:
  - Server, Storage Provisioning – No instance admin access
  - Patching
  - Backup
  - Recovery
  - Failure Detection and Repair
  - Synchronous Replication with Multi-AZ high availability
  - Read Replicas

# Benefits

- Lower Administrative Burden
- Performance
- Scalability
- Availability and Durability
- Security
- Manageability
- Cost Effective
- Choice of Instance Classes, Storage Options

# RDS Concepts/Terminology

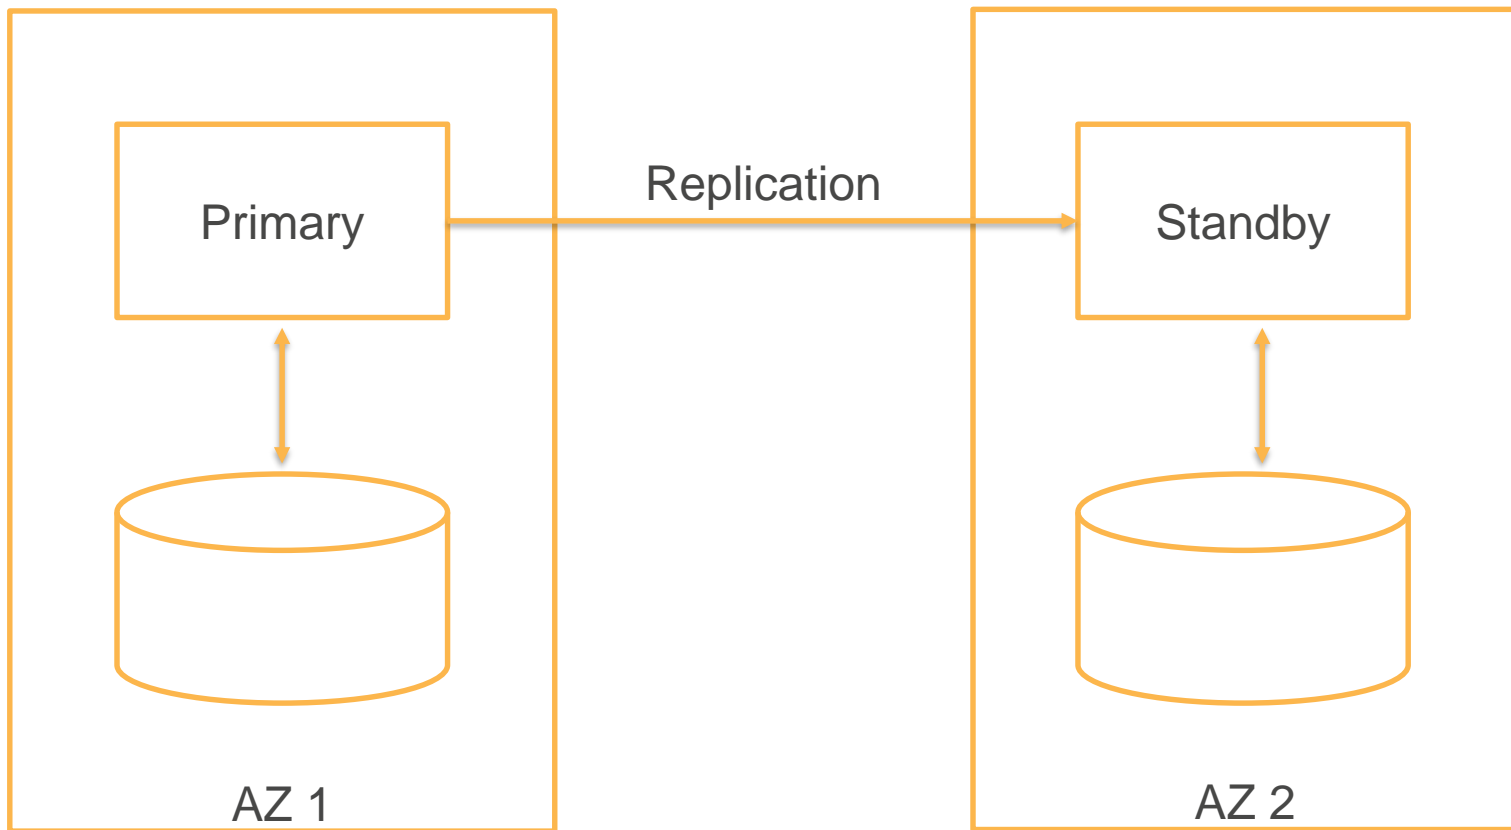
- DB Instance – A database server with your choice of DB Engine.
- EBS Storage - 5GB to 6 TB (depends on DB engine)
- Databases – Each DB Instance contains one or more customer defined databases
- VPC – Configure DB Instance to run inside your VPC
- Multi-AZ Deployment – Maintains Primary and Stand-by instances in two separate AZs with synchronous replication

# RDS Concepts

- [Primary](#) Instance – Handles all application traffic for read/write
- [Standby](#) Instance – Synchronous copy of Primary used for redundancy, failover support, system backups
- Security Group – Control access to DB Instance
- DB Parameter Group – Customize DB Engine and apply to one or more DB instances of same type
- DB Option Group – Optional features – memcache with MySQL, Data Encryption with SQL Server



# Relational Database Engines – High Availability



# Aurora

## Overview

- MySQL and PostgreSQL compatible
- Performance: 5X MySQL, 3X PostgreSQL
- Six way data replication across 3 Azs
  - For Writes, four copies must be stored safely before transaction is successful
- In case of primary crash, a read replica is promoted as primary – typically under 60 seconds
- Low latency read replicas (up to 15)
- Support for Cross Region Replication

# Aurora

- Cluster Endpoint

- Points to Current Primary Instance
- Suitable for Writes and Reads

mydbcluster.cluster-123456789012.us-east-1.rds.amazonaws.com:3306

- Reader Endpoint

- Points to Read Replicas
- Suitable for Reads
- Multiple Read Replicas are load balanced at connection level

mydbcluster.cluster-ro-123456789012.us-east-1.rds.amazonaws.com:3306

- Instance Endpoint

- Points to Individual Aurora Instance

# Aurora Serverless (GA as of Aug 2018)

- [Aurora Serverless](#) - Suitable for use cases that are intermittent or unpredictable
- Specify Minimum, Maximum Aurora Capacity Units (ACU)
- [1 ACU](#) is ~2 GB of Memory with corresponding CPU/Network
- [Pricing](#) 1 ACU is \$0.06 per hour + Storage + I/O
- Aurora Serverless automatically scales up and down based on load
- [Scaling](#) is rapid – uses a pool of warm resources

# Aurora Serverless

- Storage and Processing are separate – scale down to zero processing and pay only for storage
- [Automatic Pause and Resume](#) – Configurable period of inactivity after which DB Cluster is Paused
  - Default is 5 minutes
  - When paused, you are charged only for Storage
  - Automatically Resumes when new database connections are requested

# Pricing

- [Instance Class](#) – On-demand, Reserved
- Storage – General, Provisioned
- Backup Storage
  - First 100% provisioned storage is free
- I/O Requests per month
- Data Transfer

# Demo – RDS VPC, Client Preparation

- Setup VPC
- Setup Subnet Group
- Setup Security Groups
- Setup EC2 instance with MySQL Client
- Setup [MySQL Workbench](#) (Visual Editor)

# Demo – MySQL Multi-AZ Deployment

- Setup MySQL Database with multi-AZ option
- Note: Multi-AZ deployment is not part of free-tier, so there will small charge. We will be using micro instances.
- Connect to Primary Instance using DNS Host Name
- Create Table, Add Data
- Demonstrate Fail-Over (primary down, standby is the new primary)
- Connect to new Primary and Query



# Demo – Read Replica

Objective: Direct Read only traffic to a replica

- Create a new Read Replica
- Connect to Primary to Add new data
- Connect to Read Replica and Query

# Demo – Backup and Snapshot

- Automated Backup
- Manual Backup
- Restoring backup to a new Instance

# Instance Lifecycle

- Maintenance and Upgrades
  - Weekly Maintenance Windows
  - RDS can automatically apply patches to DB Instance or OS
  - Manually apply pending patches
  - Defer maintenance items
  - Maintenance Items marked as “Required” cannot be deferred indefinitely
  - Multi-AZ deployment: Maintenance is performed on standby, promote to primary and perform maintenance on old primary

# Modifying Instance

Modifying Existing DB Instance – What can you change?

- Apply Immediately – to apply changes right away
- Defer to next scheduled maintenance window

# Security

- [IAM](#) Control for RDS Resources

Note: You cannot connect to DB Instance with IAM Users.  
DB requires separate account

- [Encryption](#)
  - Data at rest
  - Automated Backup
  - Snapshot
  - Read Replicas
- [SSL](#) Connection

# Customization

Option Groups – Manage optional features

Examples: SQL Server Enable DB Engine level encryption for data files, Memcached support for MySQL

DB Parameter Groups – Manage Engine Configuration

RDS Provides a default configuration based on DB Engine, Instance Class, Allocated Storage

# Monitoring

## CloudWatch Metrics for RDS Instances

### Enhanced Monitoring

Enhanced Monitoring Agent running on DB Instance measures metrics like CPU

CloudWatch metrics gathers CPU metrics from hypervisor

# Architectural Best Practices

[AWS Summit Series 2016 | Chicago - Big Data Architectural Patterns and Best Practices on AWS](#)

[AWS re:Invent 2016: ElastiCache Deep Dive: Best Practices and Usage Patterns \(DAT306\)](#)



# AWS Programmatic Integration

# AWS Application Integration

- Setup Python Environment
- Install Boto3 AWS SDK
- Interact with AWS using Python

[SDKs in several supported languages](#)

# Demo – Python Boto3

Objective: Demonstrate how to programmatically connect to AWS services

1. Establish session with user credentials – same credential that was configured with Command Line interface
2. Query S3 for list of buckets
3. In DynamoDB Demo, we will use Boto3 to manage Tables, Load, Query Data

# DynamoDB

# DynamoDB

- Fully Managed NoSQL Database Service
- [Partitions](#) - Store and Retrieve any amount of data
- Automatic deletion of expired items – Time To Live
- Automatic replication of data across all AZs in a region
- Automatic scaling based on customer specified Read and Write Capacity
- [Read Consistency](#) (Eventual, Strong)
- Downloadable Developer Version

# DynamoDB Core Concepts

## Tables

- Items

- Attributes

## Primary Key

- Partition Key

- Partition Key and Sort Key

## Secondary Indexes

- Global Secondary Index

- Local Secondary Index

# DynamoDB Core Concepts

## DynamoDB Streams

- Captures data modification events in DynamoDB Tables
- Ordered Set of events
- Near real-time
- Lifetime of 24 hours
- [Figure: Lambda functions to process stream events](#)

Examples: New customer – welcome email, Add new product to ElastiCache or ElasticSearch

# Provisioned Throughput

- Consistent Low Latency Performance
- [Read Capacity](#) Units
- [Write Capacity](#) Units
- Modify any time
- Reduce cost – [Purchase Reserved Capacity](#)



# Pricing

## DynamoDB Pricing

- Provisioned Read/Write Capacity
- Index Storage
- DynamoDB Streams
- Capacity Reservation (lower cost with 1 year commitment)
- Data Transfer

# Demo – Movie Database

- DynamoDB interaction using Python
- Create Tables, Load, Query Data
- Movie demos in other programming languages are available [here](#)
- [Sample Movie Data](#)

# Monitoring

## CloudWatch Metrics

Access from Table Metrics or through CloudWatch